webfold WARC/S3 benchmark surface

One measured acquisition mode inside a live URL, CDP, DOM, document, archive, and object-store layer.

.seed parallel carry path

26,675 p/s

6.6 s wall; 3.17 s processing

4.95 GB seed; 837 MB Markdown

84,615 pages; 99.8% coverage

raw WARC parallel

14,238 p/s

16.9 s wall

15.9 GB full corpus resident

84,749 pages; 1,876 MB Markdown

single-thread stream

2,943 p/s

29 s wall

35 MB RSS

O(1) memory seed path

S3 streaming

1,169 p/s

syncfold 8-worker

73 s wall; 130 MB RSS

network-bound at about 67 MB/s

Daft full pipeline

668 p/s

Python UDF extraction/render

webfold ratio: 39.9x

Daft L1 parser alone: 13,494 p/s

SIMD scanner

13,978 p/s

hashtrinity explicit SIMD

naive: 13,051 p/s

memchr reference: 14,238 p/s

.seed scaling

9.1x

8-thread carry vs 1-thread stream

buffer approach: 15,745 p/s

carry approach avoids 45.6 GB memcpy

transport comparison

1,077 p/s

rclone FUSE -> webfold

rclone cat -> webfold: 941 p/s

syncfold serial: 546 p/s

corpus and output

15.9 GB

254,251 WARC records

84,749 HTTP responses

seed ratio 3.15:1; 15,154 blocks

Interpretation local WARC extraction is CPU/memory-path dominated · S3 runs are network-bound · browser/document surfaces require separate measurements